

## Boosting

CSE 250B

## Weak learners

It is often easy to come up with a **weak classifier**, one that is only slightly better than random guessing.

$$\Pr(h(X) \neq Y) = \frac{1}{2} - \epsilon$$

A learning algorithm that can consistently generate such classifiers is called a **weak learner**.

Is it possible to systematically boost the quality of a weak learner?

Blueprint:

- Think of the data set  $(x^{(1)}, y^{(1)}), \dots, (x^{(n)}, y^{(n)})$  as a distribution  $D$  on  $\mathcal{X} \times \mathcal{Y}$
- Repeat for  $t = 1, 2, \dots$ :
  - Feed  $D$  to the weak learner, get back a weak classifier  $h_t$
  - Reweight  $D$  to put more emphasis on points that  $h_t$  gets wrong
- Combine all these  $h_t$ 's somehow

## Adaboost

Assume  $\mathcal{Y} = \{-1, 1\}$ . Given:  $(x^{(1)}, y^{(1)}), \dots, (x^{(n)}, y^{(n)}) \in \mathcal{X} \times \mathcal{Y}$ .

- 1 Initialize  $D_1(i) = 1/n$  for all  $i = 1, 2, \dots, n$
- 2 For  $t = 1, 2, \dots, T$ :
  - Give  $D_t$  to weak learner, get back some  $h_t : \mathcal{X} \rightarrow [-1, 1]$
  - Update distribution:

$$r_t = \sum_{i=1}^n D_t(i) y^{(i)} h_t(x^{(i)}) \in [-1, 1] \quad (h_t \text{'s margin of success})$$

$$\alpha_t = \frac{1}{2} \ln \frac{1 + r_t}{1 - r_t} \in \mathbb{R} \quad (\text{strength of update})$$

$$D_{t+1}(i) \propto D_t(i) \exp(-\alpha_t y^{(i)} h_t(x^{(i)}))$$

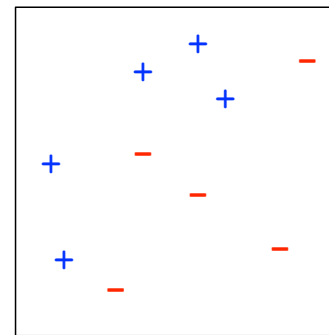
- 3 Final classifier:

$$H(x) = \text{sign} \left( \sum_{t=1}^T \alpha_t h_t(x) \right)$$

## Example

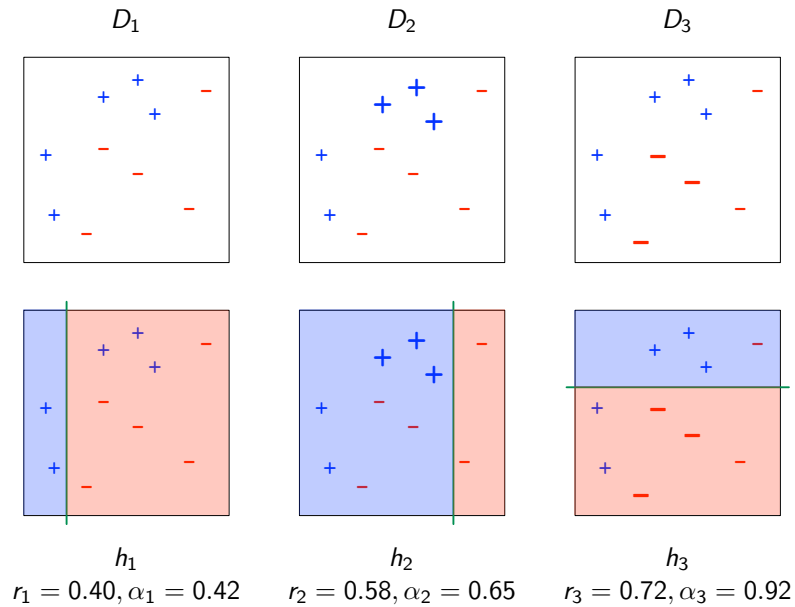
(From Freund and Schapire's tutorial.)

Training set:

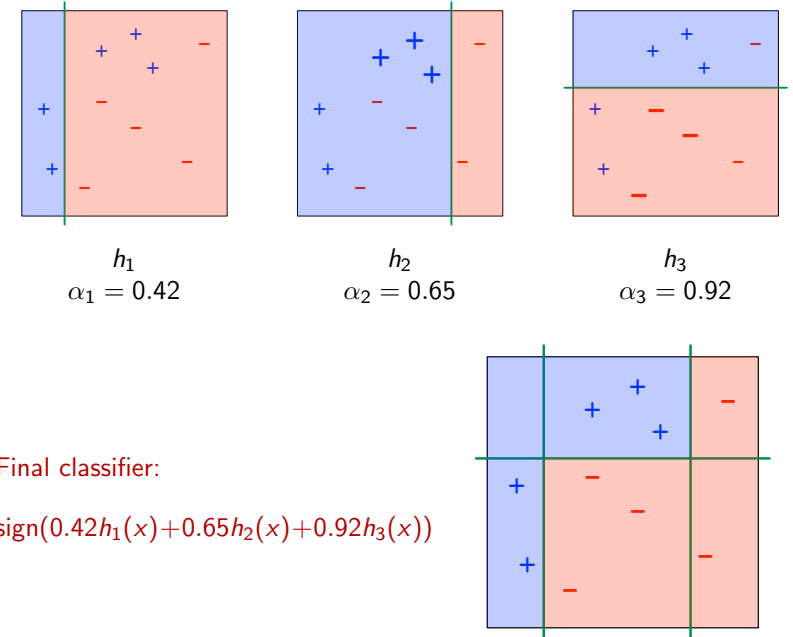


Weak classifiers: single-coordinate thresholds, popularly known as “decision stumps” (in this case, horizontal and vertical lines)

## Example, cont'd

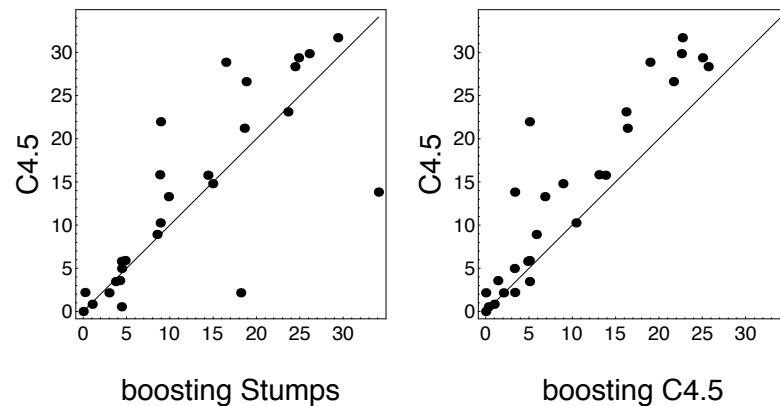


## Example, cont'd



## Boosting versus decision tree

Freund and Schapire: results on 27 benchmark data sets from the UCI repository.



## Training error dropoff

The surprising power of a weak learner:

**Theorem.** Suppose that on each round  $t$ , the weak learner returns a classifier  $h_t$  that differs from random by some margin  $\gamma$ :

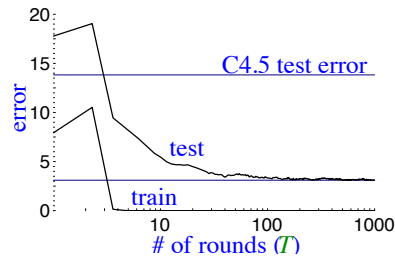
$$\left| \sum_{i=1}^n D_t(i) y^{(i)} h_t(x^{(i)}) \right| \geq \gamma.$$

Then after  $T$  rounds the training error is at most  $e^{-\gamma^2 T/2}$ .

Presumably, there will come a time  $T$  at which further reductions in training error are simply overfitting and will cause test error to rise?

## Overfitting?

Freund and Schapire: boosting decision trees for “letter” dataset.

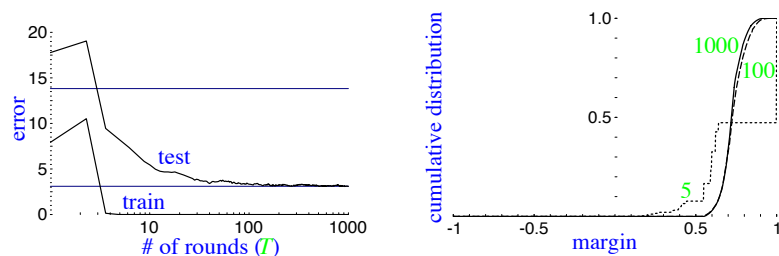


- Test error does not increase, even after 1000 rounds (total size over 2 million nodes)
- Test error keeps dropping even after training error has gone to zero:

	# rounds		
	5	100	1000
train error	0.0	0.0	0.0
test error	8.4	3.3	3.1

## Example revisited

Look at cumulative distribution of margins of training examples:



	# rounds		
	5	100	1000
train error	0.0	0.0	0.0
test error	8.4	3.3	3.1
% margins $\leq 0.5$	7.7	0.0	0.0
minimum margin	0.14	0.52	0.55

## Looking at the margin

Recall the final classifier (with weights normalized to sum to 1):

$$H(x) = \text{sign} \left( \underbrace{\frac{\sum_t \alpha_t h_t(x)}{\sum_t |\alpha_t|}}_{\text{call this } f(x)} \right)$$

The **margin** of this classifier on a particular  $(x, y) \in \mathcal{X} \times \{-1, 1\}$  is:

$$(\text{fraction of votes correct}) - (\text{fraction incorrect}) = yf(x) \in [-1, 1].$$

Did  $H$  just barely get it right? Or definitively?

- Intuitively: the larger a classifier’s margins on the training data, the better its generalization – that is, the lower its true error.
- There are mathematical results that make this precise.
- Adaboost seems to increase the margins on the training points even after training error has gone to zero.

## Another view of boosting

Let  $\mathcal{H}$  denote the set of base classifiers  $\mathcal{X} \rightarrow \{-1, 1\}$ .

For instance,  $\mathcal{H} = \{\text{decision stumps}\}$ .

Imagine a representation of  $x \in \mathcal{X}$  in which each  $h \in \mathcal{H}$  corresponds to a feature:

$$\Phi(x) = (h(x) : h \in \mathcal{H})$$

The final classifier found by boosting,

$$H(x) = \text{sign} \left( \underbrace{\sum_t \alpha_t h_t(x)}_{\text{call this } f(x)} \right).$$

is a linear classifier in this enhanced space.

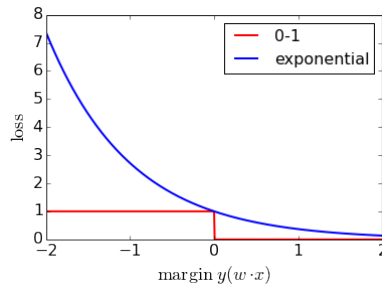
What kind of linear classifier does boosting return? Is it optimizing some loss function?

## Minimizing exponential loss

It can be shown that boosting is looking for the linear classifier that minimizes the **exponential loss**:

$$\frac{1}{n} \sum_{i=1}^n e^{-y^{(i)} f(x^{(i)})}.$$

This loss function is yet another convex upper bound on 0-1 loss:



Boosting is a **coordinate descent** procedure for minimizing the loss.

## Bagging

Given a data set  $S$  of  $n$  labeled points:

- For  $t = 1$  to  $T$ :
  - Choose  $n'$  points randomly, with replacement, from  $S$ .
  - Fit a classifier  $h_t$  to these points.

(For instance:  $T = 100$  or  $1000$ ,  $n' = n$ .)

Final predictor: majority vote of  $h_1, \dots, h_T$ .

The resampling and averaging reduces overfitting.

## Other ensemble methods

- 1 Endless variants of boosting
- 2 Bagging (bootstrapping aggregation)
- 3 Random forests

## Random forests

Rather like bagging decision trees, but with an additional source of randomization.

Given a data set  $S$  of  $n$  labeled points:

- For  $t = 1$  to  $T$ :
  - Choose  $n'$  points randomly, with replacement, from  $S$ .
  - Fit a decision tree  $h_t$  to these points. When building the tree, at each node restrict the split direction to be one of  $k$  features chosen at random.

(For instance:  $k = \sqrt{p}$  when data is  $p$ -dimensional.)

Final predictor: majority vote of  $h_1, \dots, h_T$ .